

Data Mining - Lab - 2

Numpy & Perform Data Exploration with Pandas

Numpy

- 1) NumPy (Numerical Python) is a powerful open-source library in Python used for numerical and scientific computing.
- 2) It provides support for large, multi-dimensional arrays and matrices, along with a collection of mathematical functions to operate on them efficiently.
- 3) NumPy is highly optimized and written in C, making it much faster than using regular Python lists for numerical operations.
- 4) It serves as the foundation for many other Python libraries in data science and machine learning, like pandas, TensorFlow, and scikit-learn.
- 5) With features like broadcasting, vectorization, and integration with C/C++ code, NumPy allows for cleaner and faster code in numerical computations.

Step 1. Import the Numpy library

```
In [2]: import numpy as np
```

Step 2. Create a 1D array of numbers

```
In [6]: a=np.arange(11)
print(a)
print(type(a))
```

```
[ 0  1  2  3  4  5  6  7  8  9 10]
<class 'numpy.ndarray'>
```

```
In [4]: a = np.arange(2,9)
a
```

```
Out[4]: array([2, 3, 4, 5, 6, 7, 8])
```

Step 3. Reshape 1D to 2D Array

```
In [7]: a = np.arange(12).reshape(3,4)
a
```

```
Out[7]: array([[ 0,  1,  2,  3],
               [ 4,  5,  6,  7],
               [ 8,  9, 10, 11]])
```

Step 4. Create a Linspace array

```
In [10]: np.linspace(0,5,20)
```

```
Out[10]: array([0.          , 0.26315789, 0.52631579, 0.78947368, 1.05263158,
                1.31578947, 1.57894737, 1.84210526, 2.10526316, 2.36842105,
                2.63157895, 2.89473684, 3.15789474, 3.42105263, 3.68421053,
                3.94736842, 4.21052632, 4.47368421, 4.73684211, 5.
                ])
```

Step 5. Create a Random Numbered Array

```
In [74]: np.random.rand(2)
```

```
Out[74]: array([0.83463274, 0.17866857])
```

```
In [75]: np.random.rand(2,4)
```

```
Out[75]: array([[0.08750821, 0.68581566, 0.14007458, 0.39346818],
                [0.31916189, 0.67722915, 0.81519024, 0.51836977]])
```

Step 6. Create a Random Integer Array

```
In [79]: np.random.randint(1,100,size=10)
```

```
Out[79]: array([58, 32, 52, 81, 10, 42,  3, 94, 45, 76])
```

```
In [78]: np.random.randint(1,100,size=(2,4))
```

```
Out[78]: array([[65, 95, 93, 71],
                [92, 83, 98, 59]])
```

Step 7. Create a 1D Array and get Max,Min,ArgMax,ArgMin

```
In [80]: arr = np.random.randint(1,100,size=10)
arr
```

```
Out[80]: array([89, 26, 63, 69,  1, 28,  7, 87, 86, 90])
```

```
In [81]: arr.max()
```

```
Out[81]: 90
```

```
In [82]: arr.min()
```

```
Out[82]: 1
```

```
In [83]: arr.argmax()
```

```
Out[83]: 9
```

```
In [84]: arr.argmin()
```

```
Out[84]: 4
```

Step 8. Indexing in 1D Array

```
In [85]: arr[8]
```

```
Out[85]: 86
```

```
In [86]: arr[1:5]
```

```
Out[86]: array([26, 63, 69,  1])
```

Step 9. Indexing in 2D Array

```
In [87]: arr2d=arr.reshape(2,5)  
arr2d
```

```
Out[87]: array([[89, 26, 63, 69,  1],  
                [28,  7, 87, 86, 90]])
```

```
In [88]: arr2d[0]
```

```
Out[88]: array([89, 26, 63, 69,  1])
```

```
In [89]: arr2d[0][1]  
arr2d[0,1]
```

```
Out[89]: 26
```

```
In [90]: arr2d[:,1,2:]
```

```
Out[90]: array([[63, 69,  1]])
```

Step 10. Conditional Selection

```
In [91]: arr[arr>4]
```

```
Out[91]: array([89, 26, 63, 69, 28,  7, 87, 86, 90])
```

```
In [92]: arr2d[arr2d>2]
```

```
Out[92]: array([89, 26, 63, 69, 28,  7, 87, 86, 90])
```

🔥 You did it! 10 exercises down — you're on fire! 🔥

Pandas

Step 1. Import the necessary libraries

```
In [1]: import pandas as pd
```

Step 2. Import the dataset from this [address](https://raw.githubusercontent.com/justmarkham/DAT8/master/datasets/users.csv)
(<https://raw.githubusercontent.com/justmarkham/DAT8/master/datasets/users.csv>)

Step 3. Assign it to a variable called `users` and use the `'user_id'` as index

```
In [2]: users = pd.read_csv('https://raw.githubusercontent.com/justmarkham/')
```

Step 4. See the first 25 entries

In [3]: `users.head(25)`

13	47	M	educator	29206
14	45	M	scientist	55106
15	49	F	educator	97301
16	21	M	entertainment	10309
17	30	M	programmer	06355
18	35	F	other	37212
19	40	M	librarian	02138
20	42	F	homemaker	95660
21	26	M	writer	30068
22	25	M	writer	40206
23	30	F	artist	48197
24	21	F	artist	94533
25	39	M	engineer	55107

Step 5. See the last 10 entries

In [4]: `users.tail(10)`

Out [4]:

	age	gender	occupation	zip_code
user_id				
934	61	M	engineer	22902
935	42	M	doctor	66221
936	24	M	other	32789
937	48	M	educator	98072
938	38	F	technician	55038
939	26	F	student	33319
940	32	M	administrator	02215
941	20	M	student	97229
942	48	F	librarian	78209
943	22	M	student	77841

Step 6. What is the number of observations in the dataset?

```
In [5]: users.shape[0]
```

```
Out[5]: 943
```

Step 7. What is the number of columns in the dataset?

```
In [6]: users.shape[1]
```

```
Out[6]: 4
```

Step 8. Print the name of all the columns.

```
In [7]: users.columns
```

```
Out[7]: Index(['age', 'gender', 'occupation', 'zip_code'], dtype='object')
```

Step 9. How is the dataset indexed?

```
In [8]: # "the index" (aka "the labels")
users.index
```

```
Out[8]: Index([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10,
               ...
               934, 935, 936, 937, 938, 939, 940, 941, 942, 943],
              dtype='int64', name='user_id', length=943)
```

Step 10. What is the data type of each column?

```
In [9]: users.dtypes
```

```
Out[9]: age          int64
gender         object
occupation     object
zip_code       object
dtype: object
```

Step 11. Print only the occupation column

```
In [10]: users['occupation']
```

```
Out[10]: user_id
1      technician
2      other
3      writer
4      technician
5      other
...
939    student
940    administrator
941    student
942    librarian
943    student
Name: occupation, Length: 943, dtype: object
```

Step 12. How many different occupations are in this dataset?

```
In [11]: users.occupation.nunique()
```

```
Out[11]: 21
```

Step 13. What is the most frequent occupation?

```
In [12]: users.occupation.value_counts().head(1)
```

```
Out[12]: occupation
student    196
Name: count, dtype: int64
```

Step 14. Summarize the DataFrame.

```
In [13]: users.describe()
```

```
Out[13]:
```

	age
count	943.000000
mean	34.051962
std	12.192740
min	7.000000
25%	25.000000
50%	31.000000
75%	43.000000
max	73.000000

Step 15. Summarize all the columns

```
In [14]: users.describe(include='all')
```

```
Out[14]:
```

	age	gender	occupation	zip_code
count	943.000000	943	943	943
unique	NaN	2	21	795
top	NaN	M	student	55414
freq	NaN	670	196	9
mean	34.051962	NaN	NaN	NaN
std	12.192740	NaN	NaN	NaN
min	7.000000	NaN	NaN	NaN
25%	25.000000	NaN	NaN	NaN
50%	31.000000	NaN	NaN	NaN
75%	43.000000	NaN	NaN	NaN
max	73.000000	NaN	NaN	NaN

Step 16. Summarize only the occupation column


```
In [15]: users['occupation'].describe()
```

```
Out[15]: count          943
         unique          21
         top      student
         freq          196
         Name: occupation, dtype: object
```

Step 17. What is the mean age of users?

```
In [16]: users.age.mean()
```

```
Out[16]: 34.05196182396607
```

Step 18. What is the age with least occurrence?

```
In [17]: users.age.value_counts().tail()
```

```
Out[17]: age
         7      1
         66     1
         11     1
         10     1
         73     1
         Name: count, dtype: int64
```

You're not just learning, you're mastering it. Keep aiming higher!



```
In [ ]:
```